

▼ Image Classification in DL

▼ Imports

```
import os
import zipfile
import pandas as pd
import cv2
import numpy as np
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from keras.preprocessing.image import ImageDataGenerator
from keras.applications import VGG16
from keras.optimizers import Adam


from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Embedding, LSTM, Conv1D, MaxPooling1D, Flatten
from PIL import Image

import matplotlib.pyplot as plt
import numpy as np
import os
import tensorflow as tf

!pip install kaggle
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Requirement already satisfied: kaggle in /usr/local/lib/python3.9/dist-packages (1.5.12)
Requirement already satisfied: certifi in /usr/local/lib/python3.9/dist-packages (from kaggle) (2022.9.24)
Requirement already satisfied: tqdm in /usr/local/lib/python3.9/dist-packages (from kaggle) (4.64.0)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.9/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.9/dist-packages (from kaggle) (2.28.1)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.9/dist-packages (from kaggle) (1.26.12)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.9/dist-packages (from kaggle) (5.0.2)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.9/dist-packages (from kaggle) (1.16.0)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.9/dist-packages (from kaggle) (1.3)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.9/dist-packages (from kaggle) (3.4)
Requirement already satisfied: charset-normalizer~=2.0.0 in /usr/local/lib/python3.9/dist-packages (from kaggle) (2.0.12)
ERROR: Operation cancelled by user
```

▼ Download and unzip the dataset

```
KAGGLE_TOKEN_JSON = None #@param {type:"raw"} KAGGLE_TOKEN_JSON:   
  
import json  
  
with open('/root/.kaggle/kaggle.json', 'w') as file:  
    json.dump(KAGGLE_TOKEN_JSON, file)  
  
!chmod 600 ~/.kaggle/kaggle.json  
  
!kaggle datasets download -d gpiosenka/cards-image-datasetclassification  
  
    cards-image-datasetclassification.zip: Skipping, found more recently modified local  
  
!unzip cards-image-datasetclassification.zip  
  
    Archive:  cards-image-datasetclassification.zip  
    replace 14card types-14-(200 X 200)-94.61.h5? [y]es, [n]o, [A]ll, [N]one, [r]ename:
```

▼ Divide into train/test/valid

```
train_dir = 'train'  
test_dir = 'test'  
validation_dir = 'valid'  
  
BATCH_SIZE = 32  
IMG_SIZE=(160,160)  
  
train_dataset = tf.keras.utils.image_dataset_from_directory(train_dir,  
                                                            shuffle=True,  
                                                            batch_size=BATCH_SIZE,  
                                                            image_size=IMG_SIZE)  
  
test_dataset = tf.keras.utils.image_dataset_from_directory(test_dir,  
                                                            shuffle=True,  
                                                            batch_size=BATCH_SIZE,  
                                                            image_size=IMG_SIZE)  
  
validation_dataset = tf.keras.utils.image_dataset_from_directory(validation_dir,  
                                                                    shuffle=True,  
                                                                    batch_size=BATCH_SIZE,  
                                                                    image_size=IMG_SIZE)  
  
    Found 7624 files belonging to 53 classes.  
    Found 265 files belonging to 53 classes.  
    Found 265 files belonging to 53 classes.
```

```
class_names = train_dataset.class_names
```

```
plt.figure(figsize=(10, 10))  
for images, labels in train_dataset.take(1):  
    for i in range(9):  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(images[i].numpy().astype("uint8"))  
        plt.title(class_names[labels[i]])  
        plt.axis("off")
```

seven of diamonds



jack of clubs



queen of spades



seven of hearts



king of spades



three of spades



four of diamonds



jack of spades



queen of hearts



```

AUTOTUNE = tf.data.AUTOTUNE

train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)
validation_dataset = validation_dataset.prefetch(buffer_size=AUTOTUNE)
test_dataset = test_dataset.prefetch(buffer_size=AUTOTUNE)

data_augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomFlip('horizontal'),
    tf.keras.layers.RandomRotation(0.2),
])

```

▼ Create Sequential Model

```

# Create a sequential model
num_cards = 52
embedding_dim = 32
max_len = 10
model = Sequential()

# Add Input layer: Conv2D
model.add(Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(160, 160

# Add Max pooling layer
model.add(MaxPooling2D(pool_size=(2, 2)))

# Add Flatten layer
model.add(Flatten())

# Add Dense layer
model.add(Dense(units=64, activation='relu'))

# Add Output layer
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

```

▼ Train the model

```

# Train the model
model.fit(train_dataset, epochs=5, validation_data=validation_dataset)

```

```

Epoch 1/5
239/239 [=====] - 200s 827ms/step - loss: -20948617216.0000
Epoch 2/5

```

```

239/239 [=====] - 185s 772ms/step - loss: -359579877376.000
Epoch 3/5
239/239 [=====] - 190s 792ms/step - loss: -1594228867072.00
Epoch 4/5
239/239 [=====] - 183s 767ms/step - loss: -4216356339712.00
Epoch 5/5
239/239 [=====] - 184s 767ms/step - loss: -8607562727424.00
<keras.callbacks.History at 0x7ff45de6c640>

```

▼ Evaluate Model

```

# Evaluate model
loss, accuracy = model.evaluate(test_dataset)
print(f'Test Loss: {loss}, Test Accuracy: {accuracy}')

9/9 [=====] - 2s 175ms/step - loss: -10685542039552.0000 -
Test Loss: -10685542039552.0, Test Accuracy: 0.01886792480945587

```

▼ CNN Architecture

```

# Create a sequential model with CNN architecture
num_cards = 52
embedding_dim = 32
max_len = 10
cnnModel = Sequential()

# Add Input layer: Conv2D
cnnModel.add(Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(160,

# Add Max pooling layer
cnnModel.add(MaxPooling2D(pool_size=(2, 2)))

# Add Flatten layer
cnnModel.add(Flatten())

# Add Dense layer
cnnModel.add(Dense(units=64, activation='relu'))

# Add Output layer
cnnModel.add(Dense(units=1, activation='sigmoid'))

# Compile the CNN model
cnnModel.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

cnnModel.summary()

```

Model: "sequential_40"

Layer (type)	Output Shape	Param #
conv2d_7 (Conv2D)	(None, 158, 158, 32)	896
max_pooling2d_7 (MaxPooling 2D)	(None, 79, 79, 32)	0
flatten_8 (Flatten)	(None, 199712)	0
dense_50 (Dense)	(None, 64)	12781632
dense_51 (Dense)	(None, 1)	65

=====
Total params: 12,782,593
Trainable params: 12,782,593
Non-trainable params: 0
=====

Train the CNN model

```
cnnModel.fit(train_dataset, epochs=5, validation_data=validation_dataset)
```

```
Epoch 1/5
239/239 [=====] - 195s 810ms/step - loss: -22373029888.0000
Epoch 2/5
239/239 [=====] - 200s 835ms/step - loss: -384290390016.000
Epoch 3/5
239/239 [=====] - 185s 774ms/step - loss: -1714013863936.00
Epoch 4/5
239/239 [=====] - 187s 778ms/step - loss: -4546518581248.00
Epoch 5/5
239/239 [=====] - 185s 773ms/step - loss: -9299767066624.00
<keras.callbacks.History at 0x7ff45dca4fd0>
```

Evaluate the CNN model on test data

```
loss, accuracy = cnnModel.evaluate(test_dataset)
```

```
print(f'Test Loss (CNN): {loss}, Test Accuracy (CNN): {accuracy}')
```

```
9/9 [=====] - 2s 175ms/step - loss: -11551248482304.0000 -
Test Loss (CNN): -11551248482304.0, Test Accuracy (CNN): 0.01886792480945587
```

▼ RNN Architecture

Preprocess images to grayscale and match expected input shape of LSTM layer

```
def preprocess_images(image, label):
```

```
    image = tf.image.rgb_to_grayscale(image) # Convert image to grayscale
```

```
    image = tf.image.resize(image, (10, 32)) # Resize image to (10, 32)
```

```
    return image, label
```

```
# Apply image preprocessing to train, test, and validation datasets
```

```

train_dataset = train_dataset.map(preprocess_images)
test_dataset = test_dataset.map(preprocess_images)
validation_dataset = validation_dataset.map(preprocess_images)

# Create RNN model
rnnModel = Sequential()

# Add LSTM layer as input layer
rnnModel.add(LSTM(units=32, activation='relu', input_shape=(max_len, embedding_dim)))

# Add Dense layer
rnnModel.add(Dense(units=64, activation='relu'))

# Add Output layer
rnnModel.add(Dense(units=1, activation='sigmoid'))

# Compile the RNN model
rnnModel.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Print model summary
rnnModel.summary()

# Add debug outputs
print("Input shape: ", rnnModel.input_shape)
print("Output shape: ", rnnModel.output_shape)
print("Number of trainable parameters: ", rnnModel.count_params())

rnnModel.fit(train_dataset, epochs=5, validation_data=validation_dataset)

Model: "sequential_52"

```

Layer (type)	Output Shape	Param #
lstm_28 (LSTM)	(None, 32)	8320
dense_77 (Dense)	(None, 64)	2112
dense_78 (Dense)	(None, 1)	65

```

=====
Total params: 10,497
Trainable params: 10,497
Non-trainable params: 0
=====
Input shape: (None, 10, 32)
Output shape: (None, 1)
Number of trainable parameters: 10497
Epoch 1/5
239/239 [=====] - 10s 34ms/step - loss: -45067787621629952.
Epoch 2/5
239/239 [=====] - 8s 31ms/step - loss: nan - accuracy: 0.01
Epoch 3/5
239/239 [=====] - 9s 37ms/step - loss: nan - accuracy: 0.01

```

```
Epoch 4/5
239/239 [=====] - 9s 37ms/step - loss: nan - accuracy: 0.01
Epoch 5/5
239/239 [=====] - 8s 31ms/step - loss: nan - accuracy: 0.01
<keras.callbacks.History at 0x7ff46f01c880>
```

Double-click (or enter) to edit

```
# Evaluate the RNN model on test data
loss, accuracy = rnnModel.evaluate(test_dataset)
print(f'Test Loss (CNN): {loss}, Test Accuracy (CNN): {accuracy}')
```

```
9/9 [=====] - 1s 10ms/step - loss: nan - accuracy: 0.0189
Test Loss (CNN): nan, Test Accuracy (CNN): 0.01886792480945587
```

▼ Pretrained Model

```
# Load pre-trained VGG16 model
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(160, 160, 3))

# Freeze base_model layers
for layer in base_model.layers:
    layer.trainable = False

from keras.applications.vgg16 import VGG16
from keras.models import Sequential
from keras.layers import Lambda, Flatten, Dense
from PIL import Image
import numpy as np
import tensorflow as tf

img_width = 160
img_height = 160

# Load the pre-trained VGG16 model
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(img_width, img_height, 3))

# Create a new model
model = Sequential()

# Add a custom Lambda layer for resizing images
model.add(Lambda(lambda x: tf.image.resize(x, (img_width, img_height))), input_shape=(None, None, 3))

# Add the pre-trained model as the base
model.add(base_model)

# Add custom layers
model.add(Flatten())
```



```
# Add Dense layer with input shape based on the output shape of the previous layer
model.add(Dense(units=512, activation='relu', input_dim=np.prod(base_model.output_shape[1
model.add(Dense(units=64, activation='relu'))

# Add Output layer
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```



Analysis

The models performed well but CNN performed well since the cards do not have a time factor. RNNs are suited for time series data. The initial sequential model was not best suited for the image data.

Analysis

The models performed well but CNN performed well since the cards do not have a time factor. RNNs are suited for time series data. The initial sequential model was not best suited for the image data.