

```
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report

# Read the Auto data
csv_pandas = pd.read_csv('auto.csv')

# Output the first few rows
print(csv_pandas.head(3))
print()

# Output the dimensions of the data
print(csv_pandas.shape)
print()

# Data exploration with code
print(csv_pandas[['mpg', 'weight', 'year']].describe())
print()

# Explore data types
print(csv_pandas.dtypes)
print()

csv_pandas['cylinders'] = csv_pandas['cylinders'].astype('category')
csv_pandas['origin'] = pd.Categorical(csv_pandas['origin'])

print(csv_pandas.dtypes)
print()

# Deal with NAs
csv_pandas.dropna(inplace=True)
print(csv_pandas.shape)
print()

# Modify columns
csv_pandas['mpg_high'] = np.where(csv_pandas['mpg'] > csv_pandas['mpg'].mean(), 1, 0)
csv_pandas.drop(columns=['mpg', 'name'], inplace=True)
print(csv_pandas.head(3))
print()

# Data exploration with graphs
sns.catplot(x='mpg_high', kind='count', data=csv_pandas)
sns.relplot(x='horsepower', y='weight', style='mpg_high', data=csv_pandas)
sns.boxplot(x='mpg_high', y='weight', data=csv_pandas)

# Train/test split
X = csv_pandas.drop(columns=['mpg_high'])
y = csv_pandas['mpg_high']
training_for_x, testing_for_x, training_for_y, testing_for_y = train_test_split(X, y, test_
```

```

print(training_for_x.shape, testing_for_x.shape)

# Logistic Regression
log_reg = LogisticRegression(solver='lbfgs', max_iter=1000)
log_reg.fit(training_for_x, training_for_y)

predicted_y = log_reg.predict(testing_for_x)

print(classification_report(testing_for_y, predicted_y))

# Decision Tree
dtc = DecisionTreeClassifier()
dtc.fit(training_for_x, training_for_y)

predicted_y = dtc.predict(testing_for_x)

print(classification_report(testing_for_y, predicted_y))

# Neural Network
neuralNetwork = MLPClassifier(hidden_layer_sizes=(16, 8), max_iter=1000)
neuralNetwork.fit(training_for_x, training_for_y)

predicted_y = neuralNetwork.predict(testing_for_x)

print(classification_report(testing_for_y, predicted_y, zero_division=0))

neuralNetwork2 = MLPClassifier(hidden_layer_sizes=(8, 4), max_iter=100)
neuralNetwork2.fit(training_for_x, training_for_y)

predicted_y = neuralNetwork2.predict(testing_for_x)

print(classification_report(testing_for_y, predicted_y))

# Analysis

# a. which algorithm performed better?
#   Algorithm 1
# b. compare accuracy, recall and precision metrics by class
#   Algorithm 1 had better metrics
# c. give your analysis of why the better-performing algorithm might have outperformed the
#   Due to more iterations, algorithm 1 outperformed
# d. write a couple of sentences comparing your experiences using R versus sklearn. Feel fr
#   Sklearn is easier than R but R is more readable

```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70.0	
1	15.0	8	350.0	165	3693	11.5	70.0	
2	18.0	8	318.0	150	3436	11.0	70.0	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite

(392, 9)

	mpg	weight	year
count	392.000000	392.000000	390.000000
mean	23.445918	2977.584184	76.010256
std	7.805007	849.402560	3.668093
min	9.000000	1613.000000	70.000000
25%	17.000000	2225.250000	73.000000
50%	22.750000	2803.500000	76.000000
75%	29.000000	3614.750000	79.000000
max	46.600000	5140.000000	82.000000

mpg	float64
cylinders	int64
displacement	float64
horsepower	int64
weight	int64
acceleration	float64
year	float64
origin	int64
name	object
dtype:	object

mpg	float64
cylinders	category
displacement	float64
horsepower	int64
weight	int64
acceleration	float64
year	float64
origin	category
name	object
dtype:	object

(389, 9)

	cylinders	displacement	horsepower	weight	acceleration	year	origin	\
0	8	307.0	130	3504	12.0	70.0	1	
1	8	350.0	165	3693	11.5	70.0	1	
2	8	318.0	150	3436	11.0	70.0	1	

	mpg_high
0	0
1	0
2	0

(311, 7) (78, 7)

	precision	recall	f1-score	support
0	0.98	0.82	0.89	50
1	0.75	0.86	0.81	20
2	0.98	0.82	0.89	50

	1	0.75	0.96	0.84	28
accuracy				0.87	78
macro avg		0.86	0.89	0.87	78
weighted avg		0.89	0.87	0.87	78

		precision	recall	f1-score	support
	0	0.94	0.92	0.93	50
	1	0.86	0.89	0.88	28

accuracy				0.91	78
macro avg		0.90	0.91	0.90	78
weighted avg		0.91	0.91	0.91	78

		precision	recall	f1-score	support
	0	0.91	0.84	0.87	50
	1	0.75	0.86	0.80	28

accuracy				0.85	78
macro avg		0.83	0.85	0.84	78
weighted avg		0.85	0.85	0.85	78

		precision	recall	f1-score	support
	0	0.00	0.00	0.00	50
	1	0.36	1.00	0.53	28

accuracy				0.36	78
macro avg		0.18	0.50	0.26	78
weighted avg		0.13	0.36	0.19	78

```

/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: Undefined
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: Undefined
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: Undefined
_warn_prf(average, modifier, msg_start, len(result))

```

